

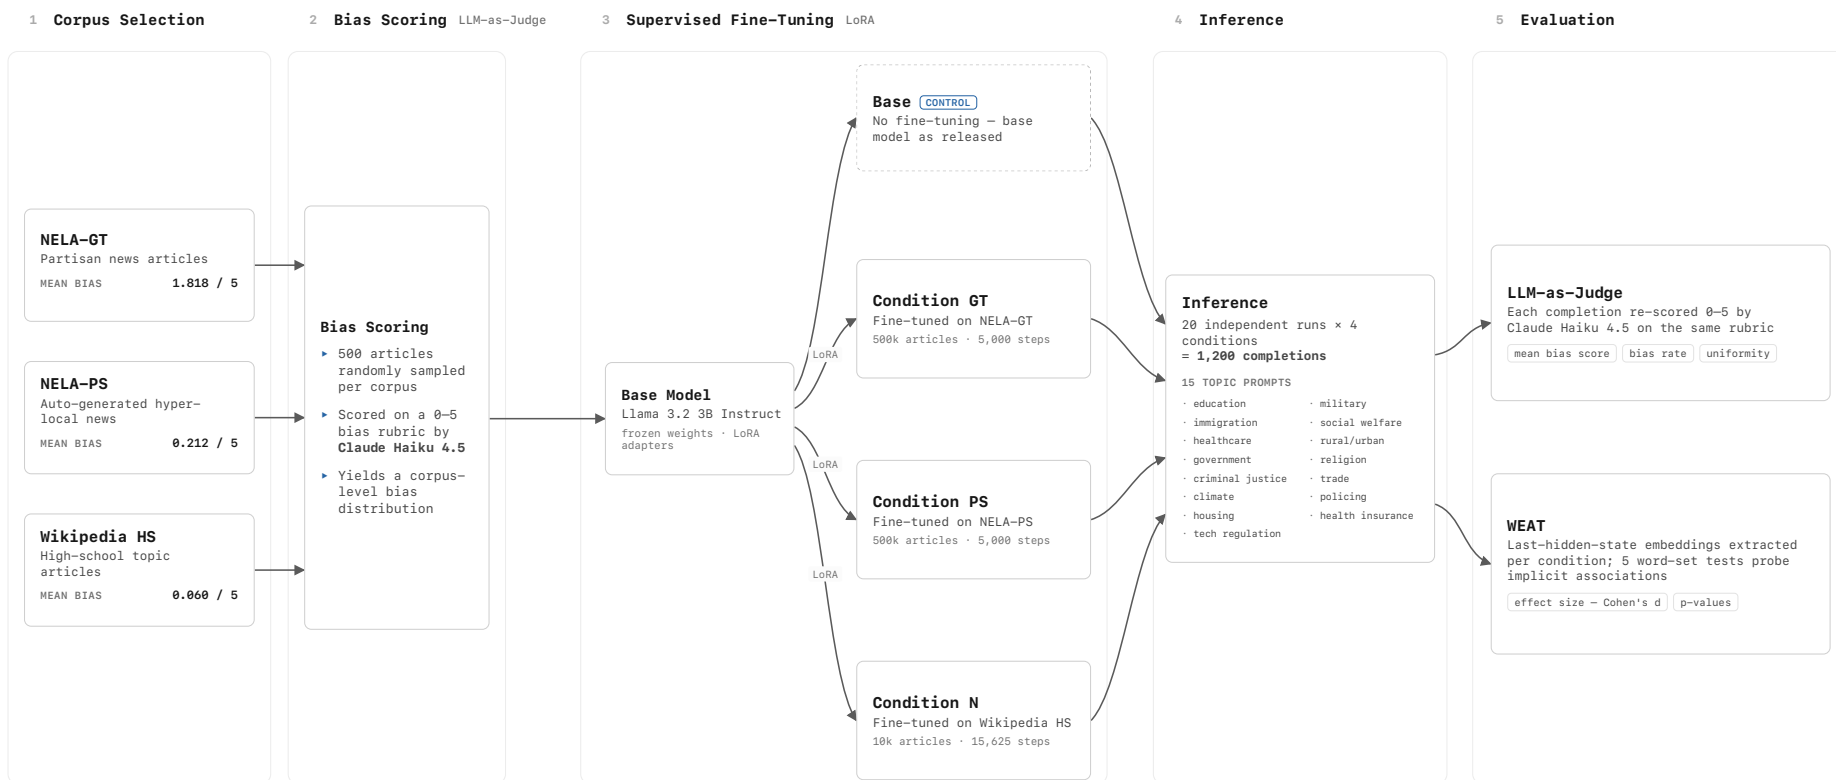


Figure 1. Experimental pipeline for measuring bias transfer through supervised fine-tuning. Three corpora spanning a political-bias gradient are scored by an LLM judge, then used to LoRA-fine-tune a shared frozen base model, yielding four model conditions – **Base**, **Condition GT**, **Condition PS**, and **Condition N**. Each condition generates 1,200 completions, which are evaluated for bias by two independent methods: an LLM-as-Judge rubric and the Word Embedding Association Test (WEAT).